


# Computer Networks

HTTP and more

Jianping Pan  
Fall 2022

## Zero RTT Connection Establishment



1. Repeat connection
2. Never talked to server before

# Last lecture

- Web, URL, HTML and HTTP basics
  - the application-layer protocol for the Web
  - following the client-server model
  - (stateless) request-reply transaction
    - HTTP request: GET / HTTP/1.0
    - HTTP response: HTTP/1.0 200 OK
  - the service expected from lower layers
    - reliable data transfer
    - normally by TCP

# Today's topics

- HTTP: advanced topics
  - fit better on TCP
    - improve HTTP efficiency
  - become stateful
    - server and client can know/remember each other
  - deal with scalability (may be left for Friday)
    - web caching and content delivery

# Web browsing examples

- `http://www.a.com/index.html`

```
<html>
```


```

```

```
<a href="http://www.a.com/y.gif">
```


```
</a>
```

```
</html>
```



- In your favorite web browser
  - URL: `http://www.a.com`
  - Q: how many HTTP requests?

# HTTP/TCP



- TCP connection establishment

- client: `connect()`; server: `accept()`

- HTTP transaction

- request: client: `send()`; server: `recv()`

- response: server: `send()`; client: `recv()`

- TCP connection release

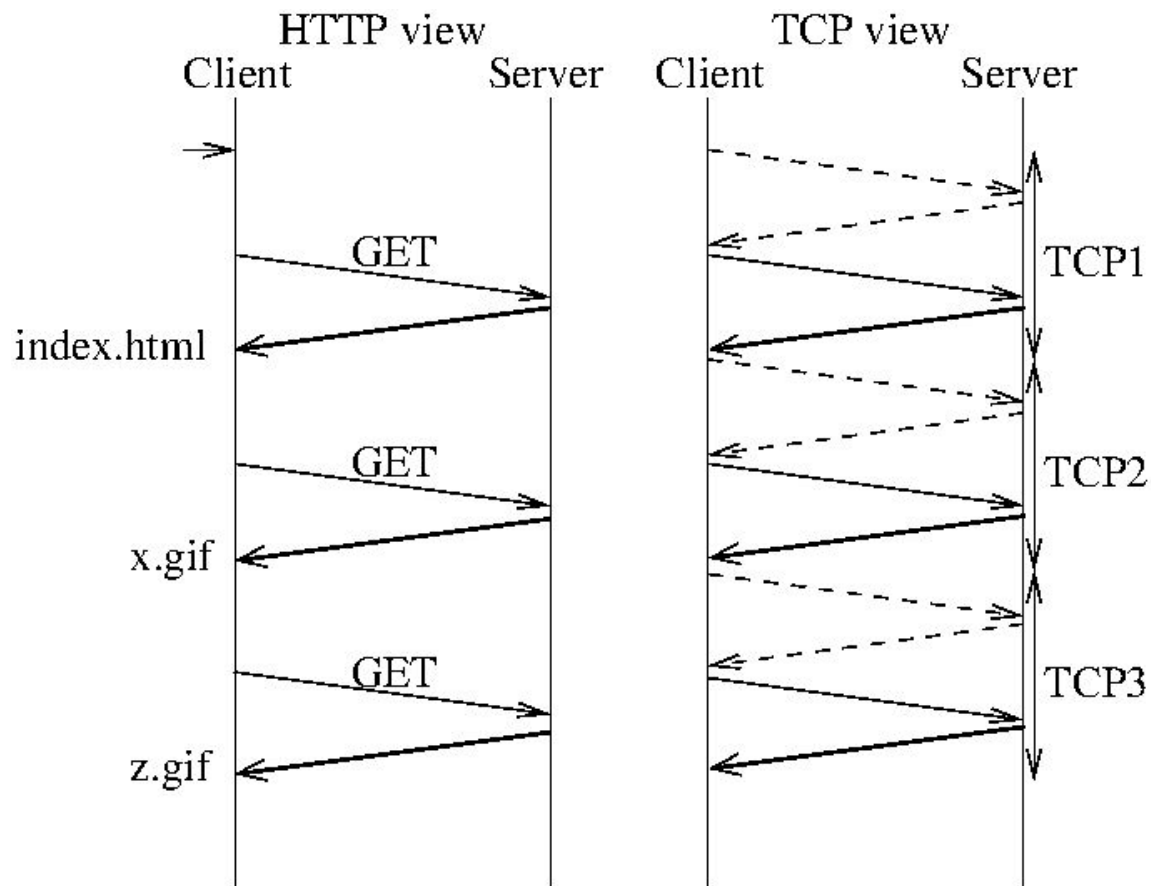
- server: `close()`; client: `close()`

Q: round-trips?


- Client is to retrieve the embedded objects

# Non-persistent HTTP

- One object per TCP connection
    - default behavior in HTTP/1.0
    - network cost:  $\sim 2 \times \text{RTT}$  per object
    - end-host cost: 1 socket() for each object
  - Performance improvement
    - parallel/concurrent connections
      - e.g., an HTML page with 2 embedded objects
- Q: cost reduced?



# Non-persistent concurrent HTTP






# Persistent HTTP


- Multiple objects through a TCP connection
  - between the **same** client/server
  - default behavior in HTTP/1.1
    - for HTTP/1.0: Connection: Keep-Alive
  - network cost:  $\sim$ RTT per object for many objects
  - to disable/close: Connection: Close
    - client/server Keep-Alive timeout
- Performance improvement: pipelining
  - $\sim 2 \times$ RTT for all objects from the same server

Q: pros and cons?


# Persistent HTTP



# Persistent HTTP + Pipelining



# Persistent + concurrent HTTP



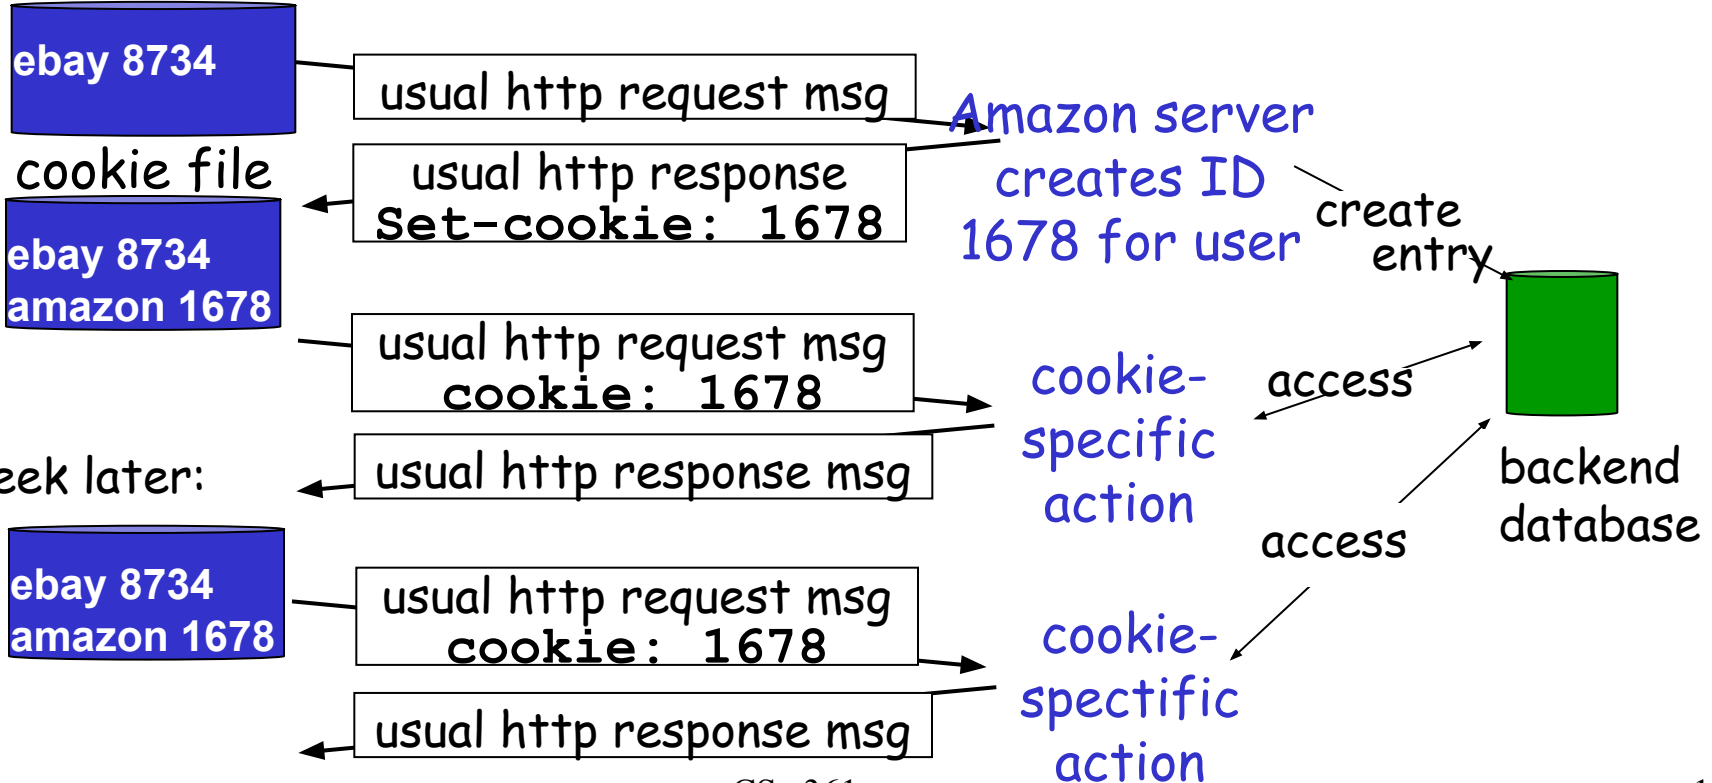
# Client-server states

- HTTP itself is stateless
  - request-reply **transactions**
- Many applications require states
  - cookie issued by server's *backend* servers
    - HTTP response header: Set-Cookie
  - client can *choose* to keep cookie
  - client presents cookie in subsequent requests
    - HTTP request header: Cookie

# HTTP cookies: an example

client


server



# Tracking client?!

- Between web servers
  - HTTP request header: Referer \*
  - Referrer!
  - referrer spamming and referrer spoofing
  - de-referring
- User security and privacy
  - e.g., cookie theft, cookie poisoning, web bug
  - browsing anonymizing

# Web caching



- Scalability issues with the client-server model
  - one server, many clients, concurrent requests
  - server load
  - network traffic
- Web caching: aggregate user requests
  - by caching responses to previous requests
  - explore locality: same requests may occur soon!
  - reduce response time and traffic load

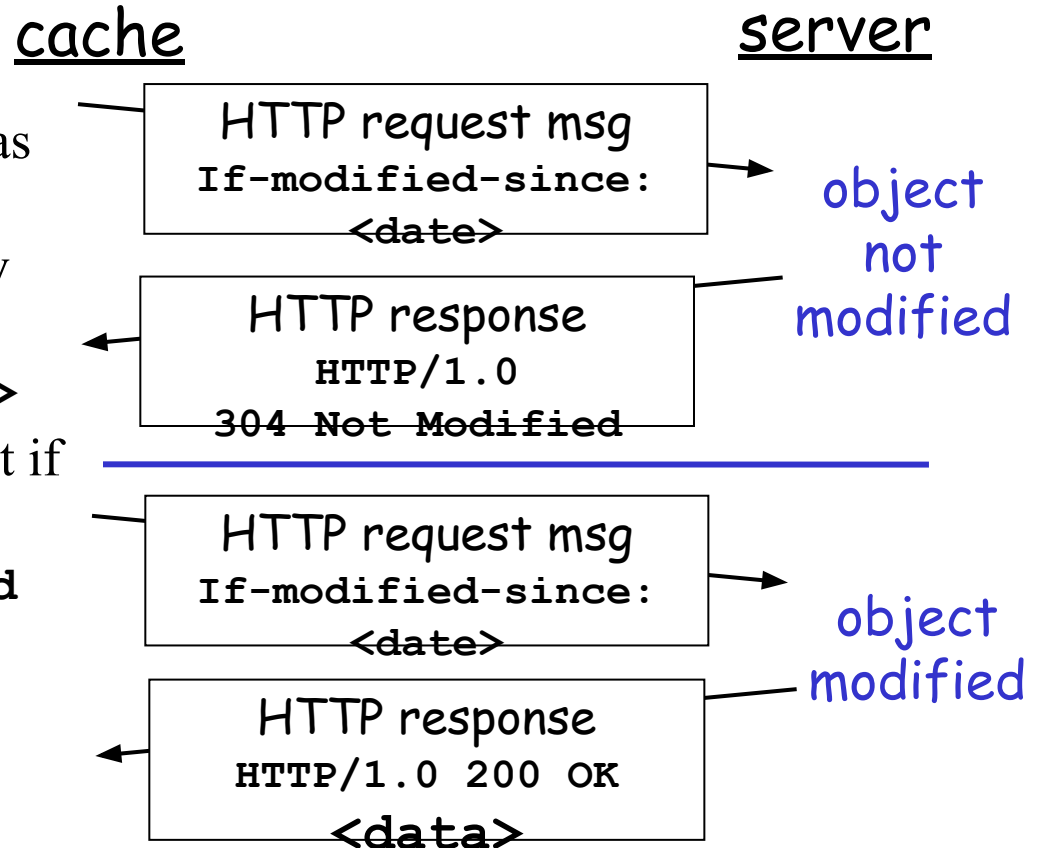


# Consistency control

- Objects retrieved from the cache
    - may be stale due to updates at origin servers
  - Strong consistency
    - HTTP request header: If-Modified-Since
    - HTTP response: HTTP/1.0 304 Not Modified
      - reduce traffic load if hit
  - Weak consistency
    - time-to-live (TTL)
      - reduce traffic load and response time if “hit”
- Q: pros and cons?

# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- server: response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# Content delivery

- Move content closer to end users
  - content distribution
- Redirect users to closer servers
  - information retrieval
  - how do they do that?
    - e.g., Akamai
    - DNS-based server selection
      - more after DNS lectures

- HTTP evolution
  - Original HTTP: only the GET method for HTML
  - HTTP/0.9 (1991): headers and additional methods
  - HTTP/1.0 (1996): performance improvement
    - **Persistent connection (Keep-Alive) becomes optional**
  - HTTP/1.1 (1997/1999): most widely used now
    - Persistent connection becomes the default
  - HTTP/2 (RFC7540, May 2015)\*
    - Based on Google's SPDY
  - HTTP/3 (RFC9000, May 2021)
- More features and how to fit better with TCP

- Google's SPDY
  - A basis for HTTP/2
  - Further reduce the page load time (PLT)
  - Prioritizing and multiplexing embedded objects
    - One connection per client to the same server
    - More flexible than HTTP pipelining (HOL blocking)
  - Encryption (TLS) and compression (Gzip)
  - Server can hint or push content/update
  - Used at Google, Facebook, Twitter, etc
- Supported in Chrome, Mozilla, Opera, etc

# This lecture

- HTTP
  - (non-)persistence, pipelining, cookies, referrers
  - web caching and content delivery
- Explore further
  - How do your favorite web browsers and servers support advanced HTTP features?
  - How's your your W1 and P1?

# Next lecture

- DNS (**KR7S2.4\***)
  - DNS server hierarchy
  - DNS resolution process
  - DNS-based server selection

## Local DNS name server

- ◊ does not strictly belong to hierarchy
- ◊ each ISP (residential ISP, company, university) has one
  - also called "default name server"
- ◊ when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date)
  - acts as proxy, forwards query into hierarchy

Slide 10/10