Computer Networks

Transmission Control Protocol

Jianping Pan
Fall 2022

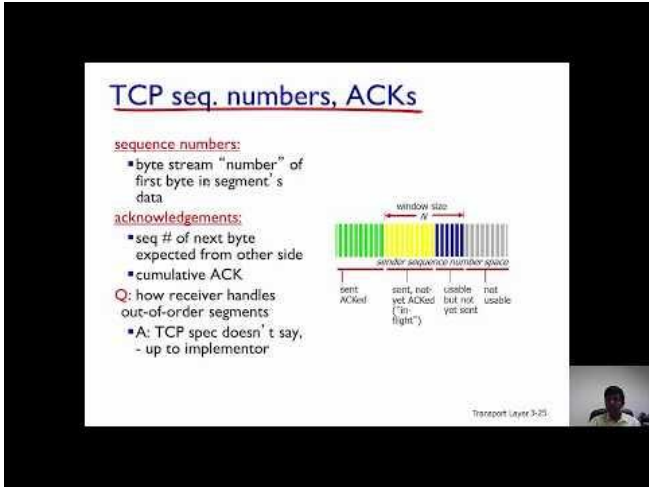

TCP

- Transmission control protocol [RFC793]
- Services provided by TCP
  - connection-oriented, point-to-point, bi-directional
  - reliable, in-sequence, stream-like
- Services provided by IP
  - packets: duplicated, lost, reordered, corrupted
- TCP protocol mechanisms
  - connection management
  - flow, error and congestion control

TCP packet header

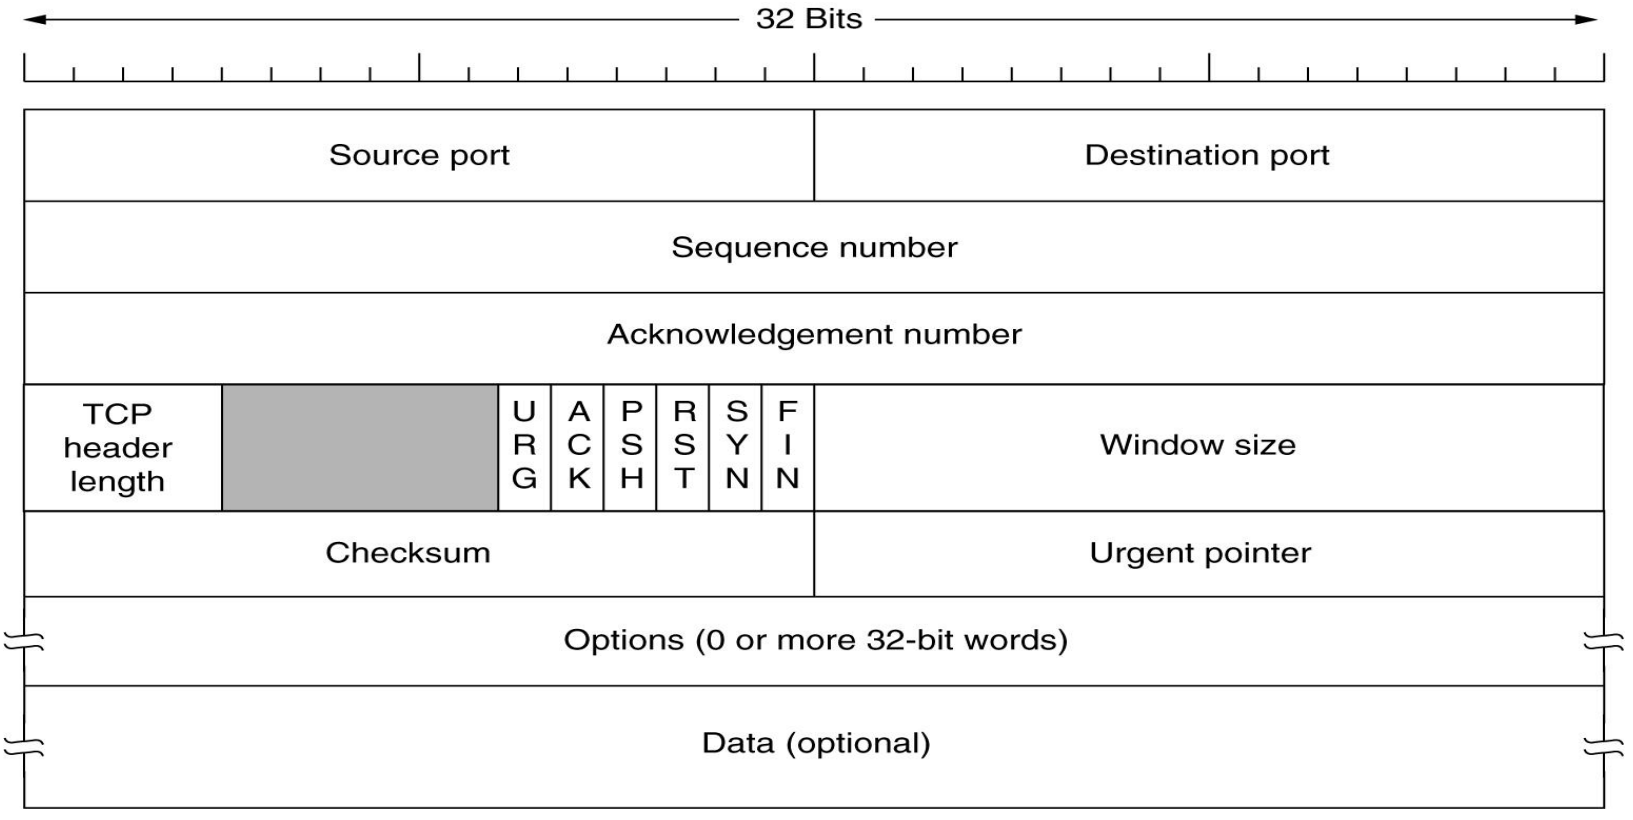

Port number

- TCP port number (16-bit)
  - source, destination port numbers
    - addressing and multiplexing
- Port number allocation (ref: iana.org)
  - well-known port numbers (0~1023, privileged)
    - e.g., 80: http; 443: https
  - registered port numbers (/etc/services)
    - http-alt 8080/tcp
  - dynamically allocated port numbers

bind(), connect()

TCP connection ID

- TCP connections
  - connection: initiator, responder
    - (initiator IP, *initiator port*, responder IP, responder port)
- One connection: one flow in each direction
  - for each flow: source, destination
    - (source IP, source port, destination IP, destination port)
  - 5-tuple (or 4-tuple when protocol ID is implied)
    - (src IP, src port, protocol ID, dst IP, dst port)
- Socket, connection, flow

listen()

connect(), accept()

send(), recv()

Q: port reuse?

Sequence number

- TCP sequence number (32-bit)
  - byte sequence for the ***first*** byte in payload
    - exception: SYN/FIN sequence number
  - random initial sequence number
    - exchanged during 3-way handshake
  - sequence number rollover

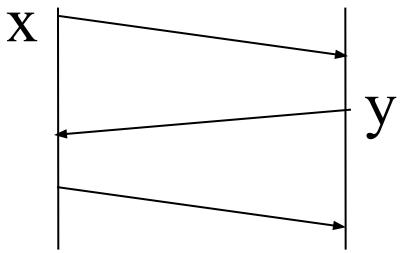

- Acknowledgment number (32-bit)
  - byte sequence for the ***next*** byte to expect

Q: why random ISN?

Sequence vs acknowledgment

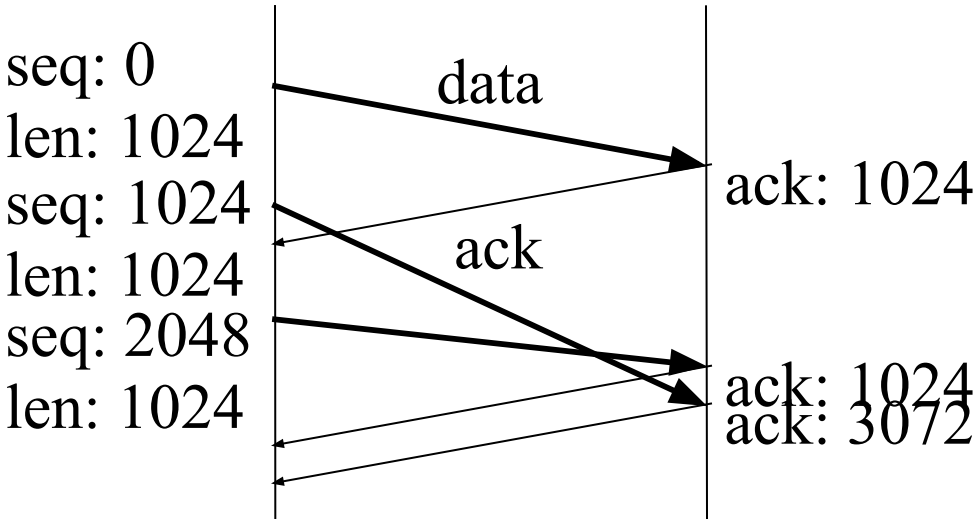

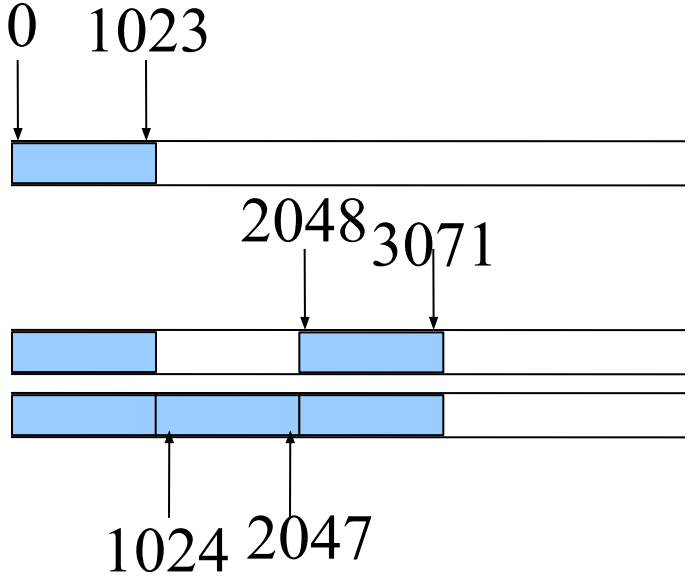

Q: byte vs packet sequence?

Header length/Data offset

- Variable-length header due to TCP options
- TCP header length/data offset (4-bit)
  - number of 32-bit words!
  - at least 5 for fixed TCP header fields *
  - maximum: 15
    - i.e., 40 bytes in total for TCP options
- TCP design feature: bit-alignment
  - fields of x-bit length at offset n*x (n: an integer)
  - 32-bit word

* why not start from 0 for 5 (or 20)?

Q: why alignment?

Control flags

- URG: urgent pointer meaningful
- ACK: acknowledgment number meaningful
- PSH: logic message boundary
- RST: connection reset
- SYN: synchronization (conn. establishment)
- FIN: finish (graceful connection release)
  - stay tuned: “TCP connection management”

Window size

- TCP window size (16-bit)
  - stay tuned: “TCP flow control”

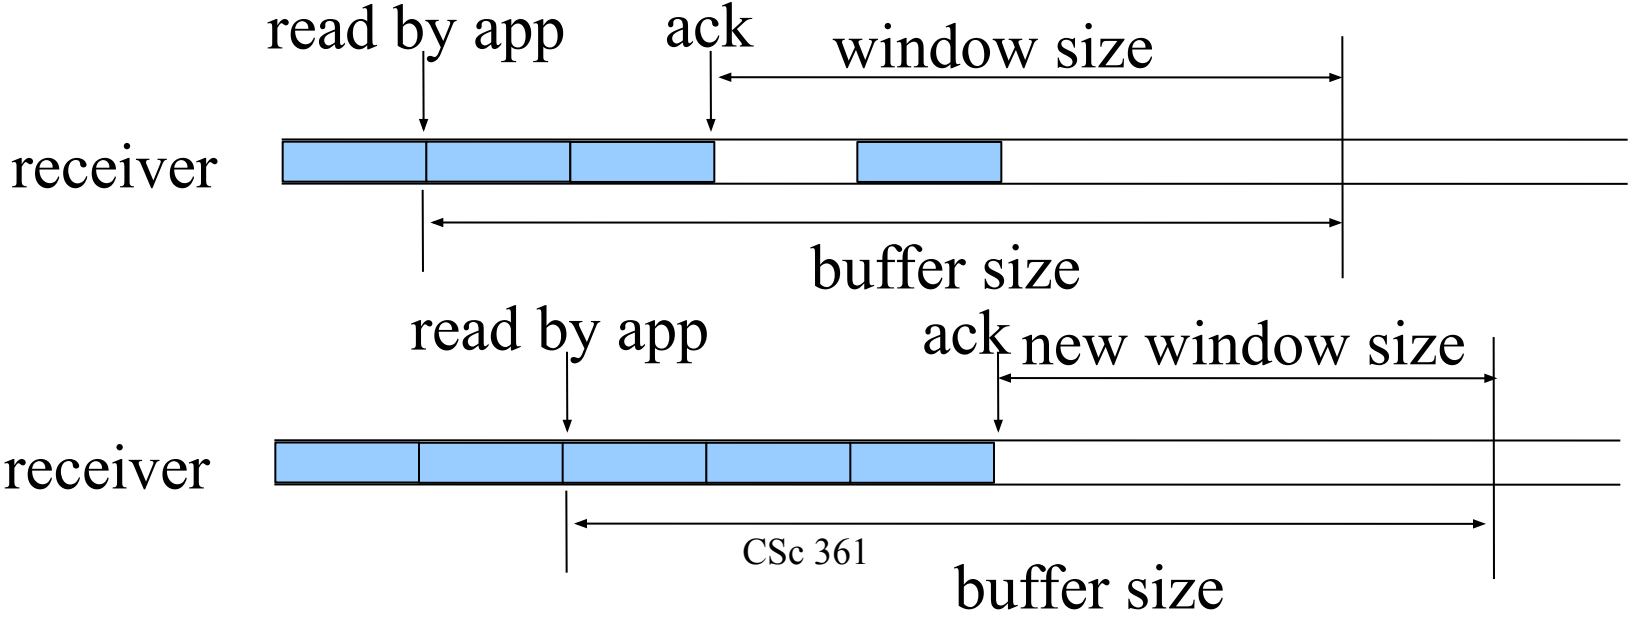

Checksum

- TCP checksum (16-bit)
  - "one's complement of one's complement sum"
    - stay tuned: "TCP error control"
  - cover
    - TCP header (including options, if any)
    - TCP payload
    - TCP pseudo header
      - source and destination IP address
      - protocol ID
      - TCP segment size

Q: why TCP pseudo header?

Stone, Jonathan, and Craig Partridge. "When the CRC and TCP checksum disagree." ACM SIGCOMM Computer Communication Review. Vol. 30. No. 4. ACM, 2000.

Urgent pointer

- TCP urgent pointer (16-bit)
  - offset of the LAST byte for urgent data
    - not (LAST+1) per RFC 1122: Host requirements
  - from the current sequence number!
  - for out-of-band (OOB) control information*
    - e.g., interrupt an ongoing file transfer
  - Socket interface
    - send(s, buf, len, MSG_OOB);
  - receiver should process the urgent data immediately

* can be overridden by MSG_OOBINLINE Q: URG vs PSH?

TCP options

- TLV-like options
  - option-kind: 1-byte
  - option-length: 1-byte, for the entire option
  - option-data:variable length
- E.g., Maximum Segment Size (MSS)
  - exchanged during connection establishment
  - default: 536 bytes
- E.g., Selective Acknowledgment (SACK)
  - stay tuned: “TCP congestion control”
- Zero-padding to keep 32-bit alignment

Q: big vs small segment?

TCP packet header (again)

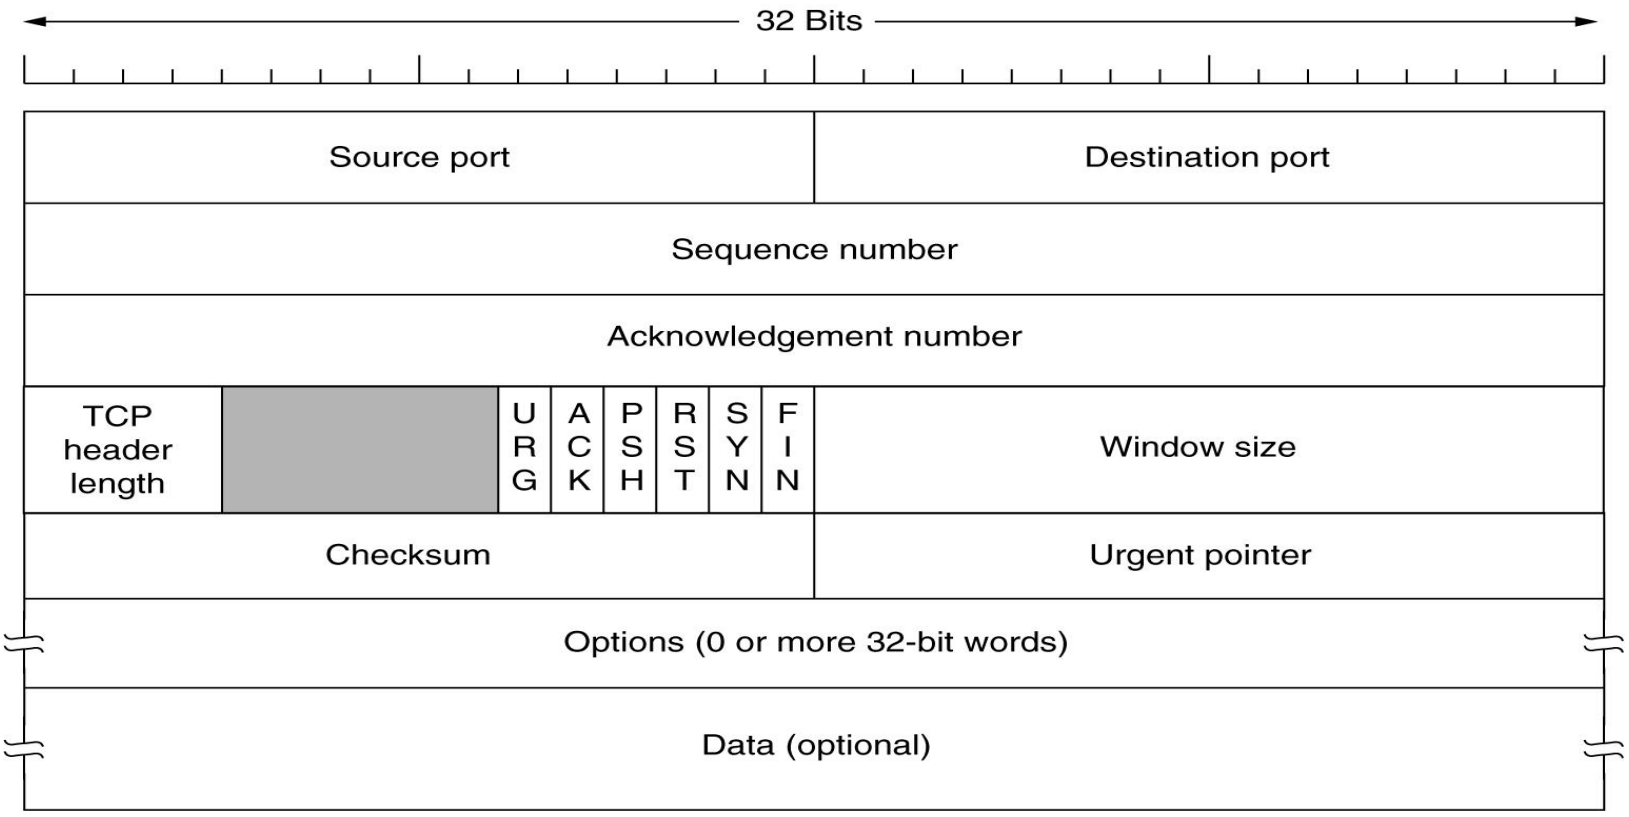

Q: any redundant fields?

This lecture

RIP, Sally!

- TCP
  - services provided by TCP
    - protocol mechanisms to support TCP services
  - TCP header fields
    - control information exchanged to support TCP
- Explore further
  - Qs on previous slides
  - the limitation of TCP and proposed changes
    - http://www.icir.org/floyd/tcp_small.html

"Research - TCP/IP over something - Air Links - Wireless Links. PAN Jianping's annotated pointers." now at http://www.cs.uvic.ca/~pan/tcpair

Next lecture

- This Friday, Oct 7, 9:30am: M1
- Next week: TCP connection management
  - read KR4: Computer Networking
    - Chapter 3 (all sections required this month)

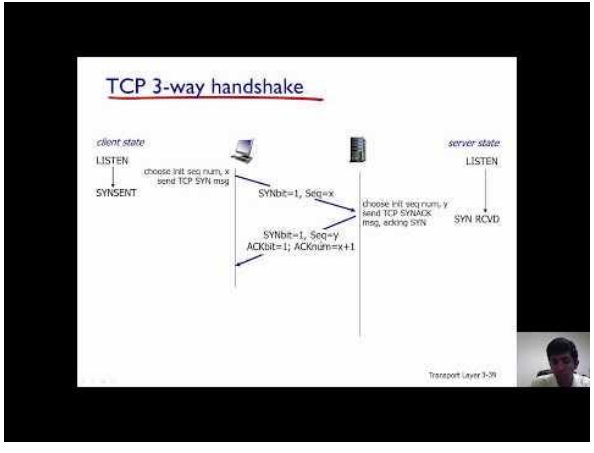

Midterm 1 (M1)

- Friday, October 7, 2022, 9:30am Victoria Time (GMT-7)
  - In ECS116 (sorry, the time/space has been determined with CAL already)
  - Let see how it goes, and whether we shall make adjustments for future ones
- Coverage
  - All lectures, tutorials and labs, and their corresponding text materials, slides, assignments, etc
  - From the beginning of the term to Wednesday, October 5, 2022, all inclusive
- Type of questions
  - Concepts questions: describe and compare related concepts
  - Role-playing questions: If I am the computer (network), what shall I do?
  - Play-a-role questions: If I was given the (real) problem to solve, what shall I do?
- Format of the exam
  - Friday 9:30am to 10:20am, Victoria time, ECS116
    - If you have CAL accommodation, the time duration (50 minutes) will be prorated at CAL
- Exams are closed book/notes/Internet, and have to be done by yourself alone