

Google @ UVic | Fall 2022

Building Your Technical Career:

- Learn how to make the most of your degree
- Prepare for a career in software engineering

Interview Prep Workshop:

- Watch a mock interview
- Solve sample problems as a group

STEP Resume Workshop (Tech):

- Learn about Google's STEP internship for 1st/2nd year students
- Get resume tips on how to best present your skills



Hosted by Google Software Engineers **Ian Sutton** and **Faesar Murad**

Google @ UVic | Fall 2022

Tuesday, October 18th

- **Building Your Technical Career:** *ECS 660* | 5pm - 6:30pm

Wednesday, October 19th

- **Office hours for 1st/2nd year students:** *ECS 223 & ECS 227* | 1pm - 3pm

Thursday, October 20th

- **Interview Prep Workshop:** *ECS 660* | 12pm - 1pm
- **STEP Resume Workshop (Tech):** *ECS 660* | 5:30pm - 6:30pm

Friday October 21st:

- **Virtual office hours:** *Online* | 12pm - 1pm

Please RSVP at <https://goo.gle/UVictoriaFall2022Events>

Check out more career-related content online at <https://careersonair.withgoogle.com/>

Computer Networks

TCP Error Control

Jianping Pan
Fall 2022

TCP round trip time, timeout

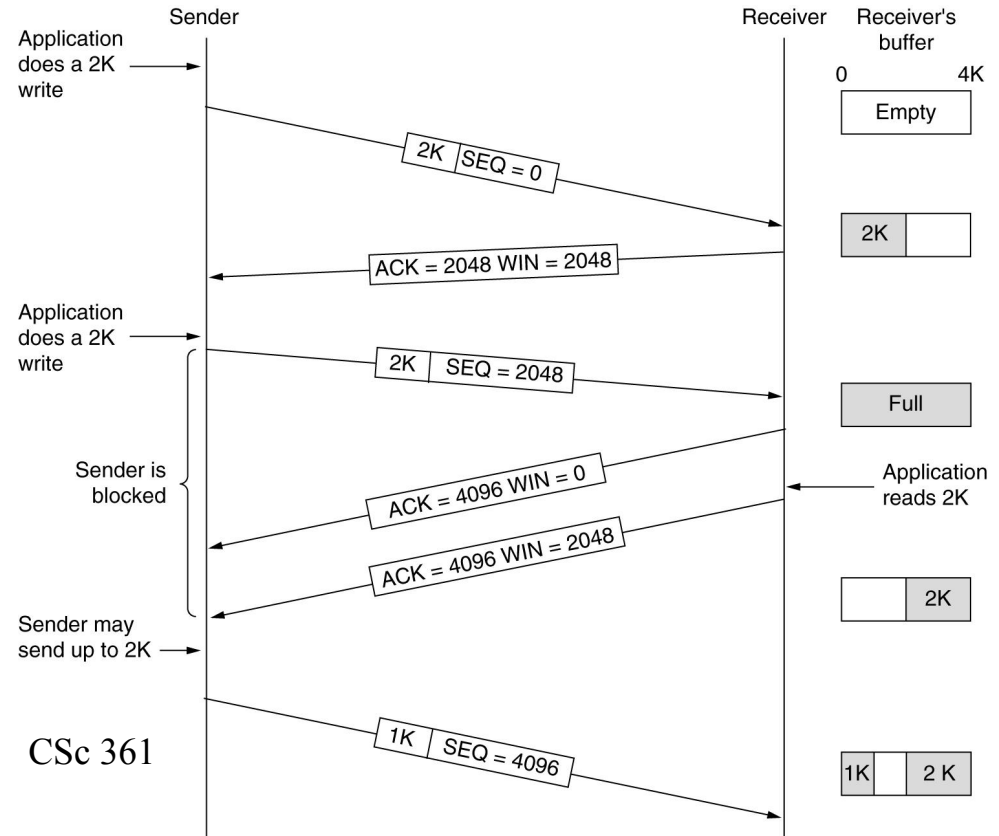
$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value: $\alpha = 0.125$



Review: TCP flow control

- Purpose
 - to avoid overflow
- Mechanism
 - sliding window
 - variable window



10/18/22

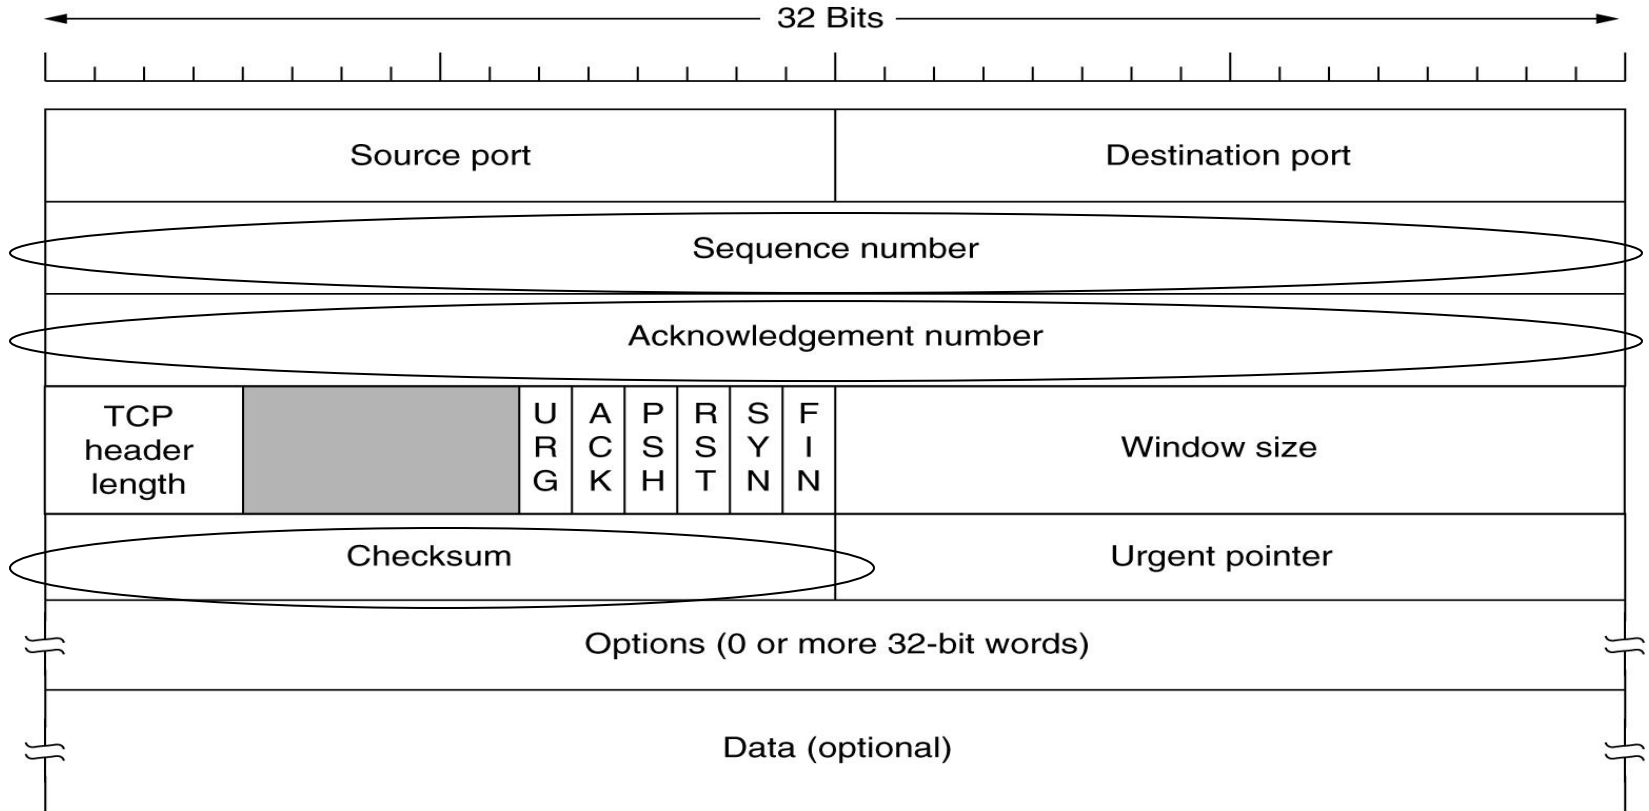
Q: seq, ack, win?

Error control

- Service provided by TCP
 - connection-oriented, **reliable** data transfer
- Service provided by IP
 - connectionless, ***unreliable*** packet delivery
 - packets may get
 - lost
 - duplicated
 - corrupted
 - reordered

Q: why?

TCP packet header



What can go wrong?

- IP packet delivery
 - lost
 - *transmission error* or network congestion
 - duplicated
 - deleted by referring to sequence number; **done**
 - corrupted
 - arrived but in “bad shape”
 - reordered
 - rearranged by referring to sequence number; **done**

Error detection

- Corrupted packets
 - detected by TCP checksum
 - action: drop!
- Lost packets
 - how do you tell if something *is* already lost?
 - TCP sender
 - timer for acknowledgment
 - TCP receiver (cumulative acknowledgment)
 - duplicate acknowledgment

TCP/IP checksum

- Algorithm: 16-bit one complement of one's complement sum with carry
 - **16-bit: padding** when necessary
 - cover: TCP header, payload, pseudo header
 - calculate: pad, sum, carry, complement=>checksum
 - verify: sum with checksum, carry, complement=>0?

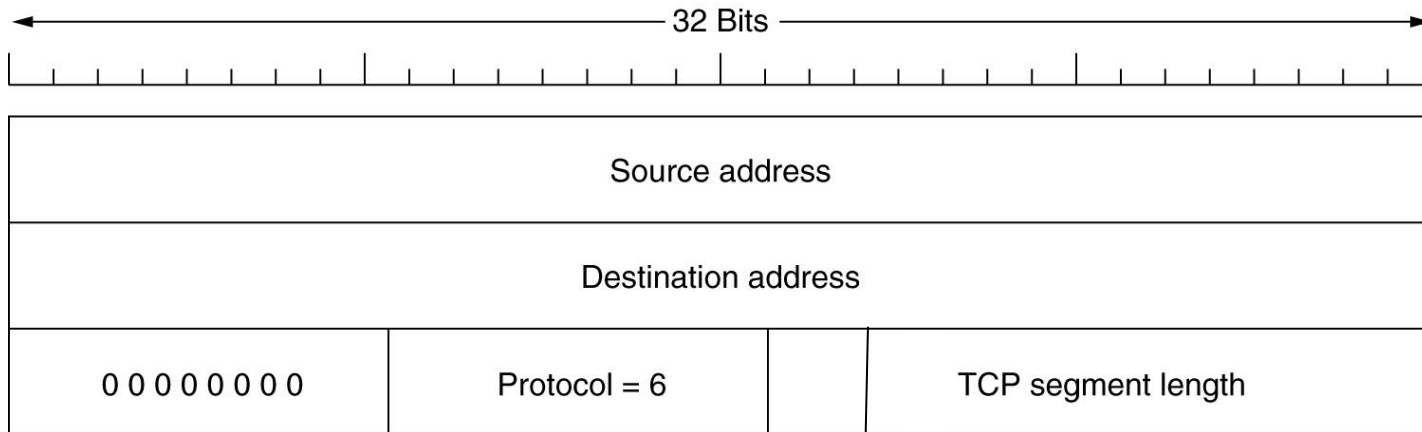
Examples

1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1

wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1	1

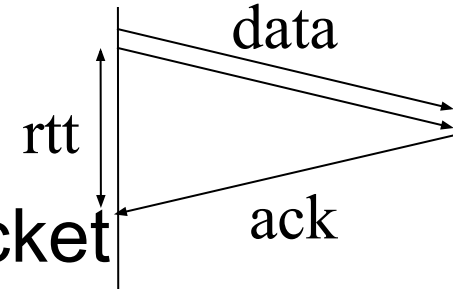
IP pseudo header

- TCP checksum also covers IP pseudo header
 - to detect mis-delivered packets by IP layer
 - include: IP addresses, protocol ID, segment length



TCP sender timer

- TCP sender
 - start a timer when sending out a packet
 - in reality: one timer per a window of packets
 - on acknowledgment “covering” this packet
 - cancel the timer and setup another one
 - if timer timeouts: *indicate* packet may be lost
- Timeout value
 - too soon: unnecessary transmission
 - too late: “slow response”

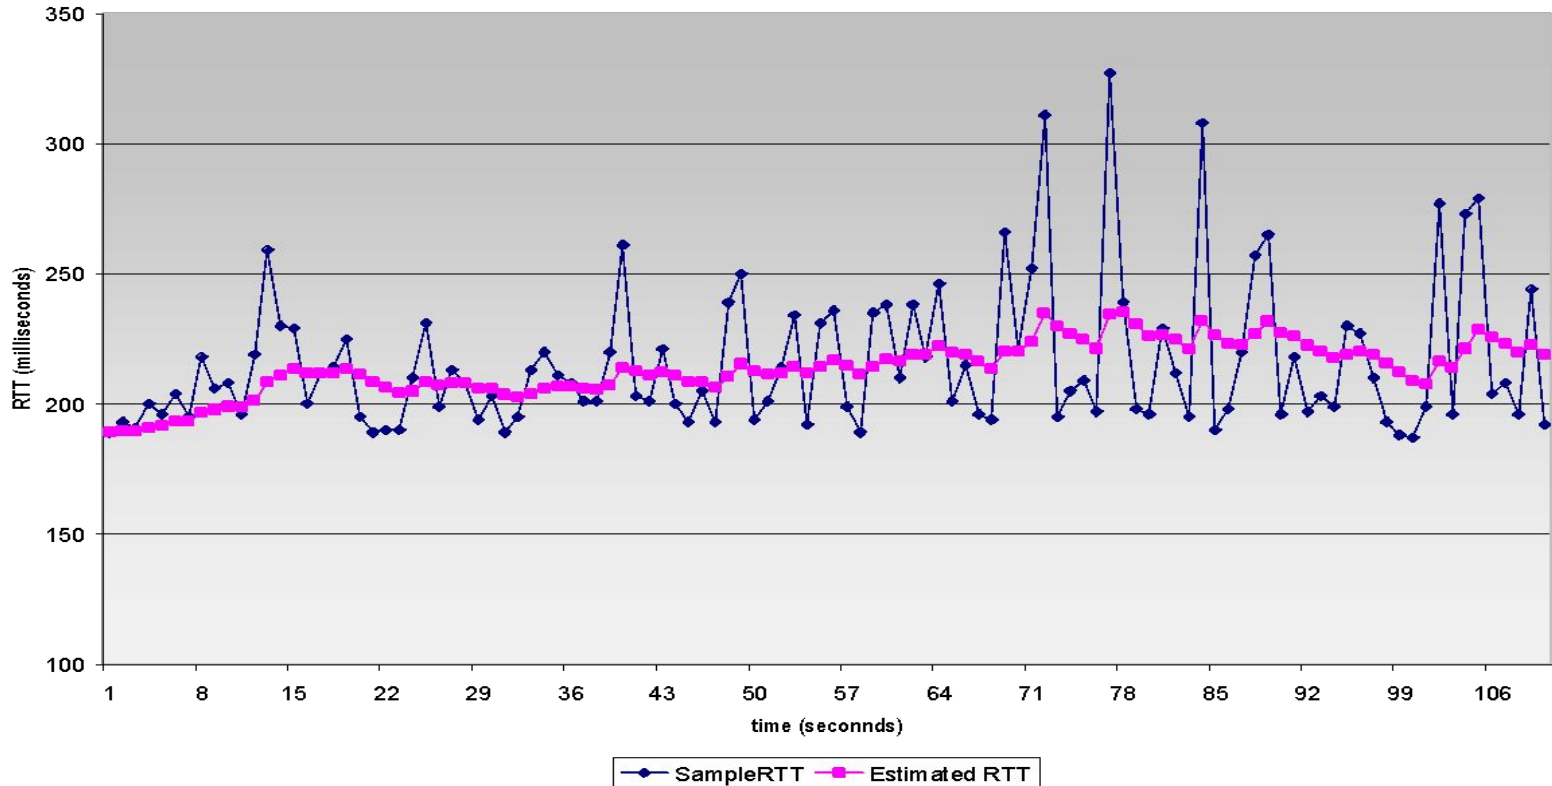


TCP round-trip time

- RTT measurement and calculation
 - RTT sample
 - time from sending a packet to receiving its ack
 - coarse-grained: 500 ms in BSD
 - ignore retransmitted packets for RTT measurement
 - smoothed RTT (SRTT)
 - exponentially weighted moving average (EWMA)
 - $SRTT_{i+1} = SRTT_i + a (RTT - SRTT_i)$
 - $a = 1/8$

EWMA example

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



TCP timeout value

- RTO calculation based on SRTT
 - RTT variance (RTTV)
 - $RTTV_{i+1} = RTTV_i + b(|RTT - SRTT_i| - RTTV_i)$
 - $b = 1/4$
 - RTO
 - $RTO = d (SRTT + c RTTV)$
 - c : initially 2, now 4
 - d : backoff factor
 - initially 1, doubled when timeout until reaching the maximum
 - initial SRTT, RTTV and minimum RTO

TCP sender (simplified with no flow control and congestion control)

send_next = InitialSeqNum

ack = InitialSeqNum

```
loop (forever) {  
    switch(event)
```

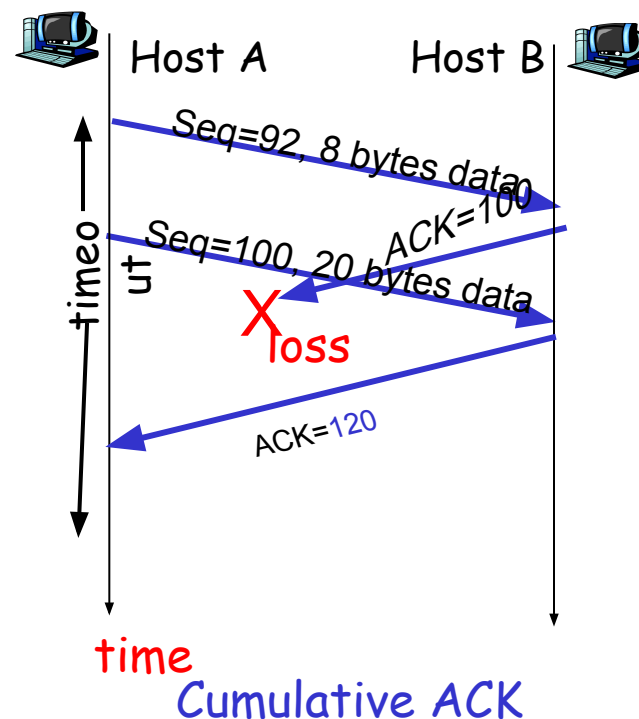
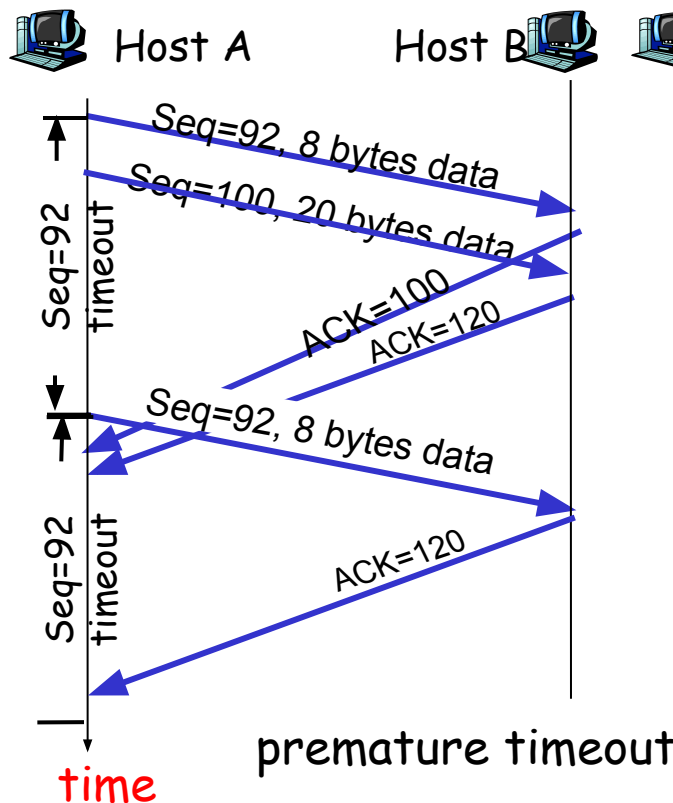
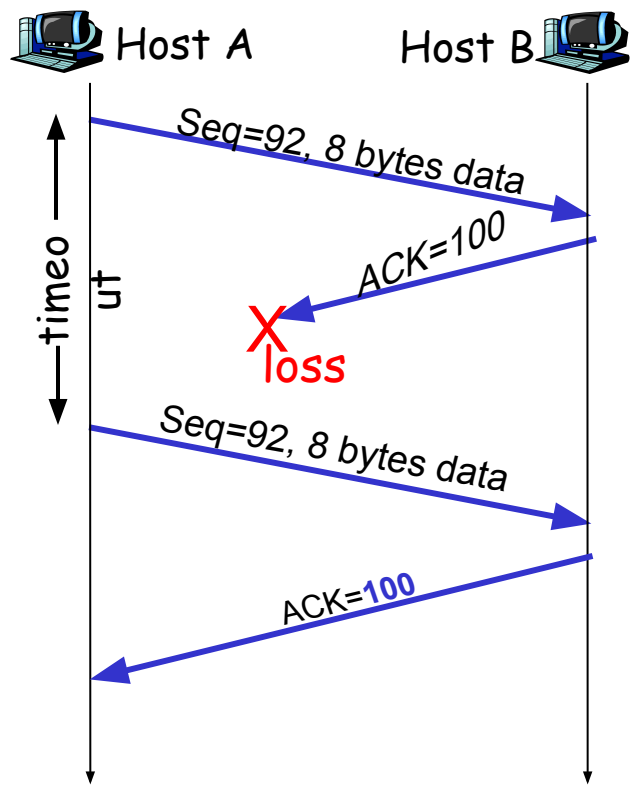
```
event: data received from application above  
    create TCP segment with sequence number send_next  
    if (timer currently not running)  
        start timer with timer's seqno = send_next  
    pass segment to IP  
    send_next = send_next + length(data)
```

```
event: timer timeout  
    retransmit not-yet-acknowledged segment with  
        smallest sequence number  
    start timer with the resent seqno
```

```
event: ACK received, with ACK field value of y  
    if (y > ack) {  
        ack = y  
        cancel timer with timer's seqno < y  
        if (timer not running && there are currently not-yet-acknowledged segments)  
            start timer  
    }
```

```
    } /* end of loop forever */
```

TCP: retransmission scenarios



lost ACK scenario

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver

TCP Receiver action

Arrival of **in-order segment** with expected seq #. **All** data up to expected seq # already **ACKed**

Delayed ACK. Wait up to **500ms** for next segment. If no next segment, send **ACK**

Arrival of in-order segment with expected seq #. **One** other segment has ACK **waiting**

Immediately send **single cumulative** ACK, ACKing both in-order segments

Arrival of **out-of-order** segment **higher-than-expect seq. #** .
Gap detected

Immediately send **duplicate ACK**, indicating seq. # of **next expected** byte

Arrival of segment that partially or completely fills gap

Immediate send **ACK**, provided that segment starts at lower end of gap
(**accept out-of-order**)

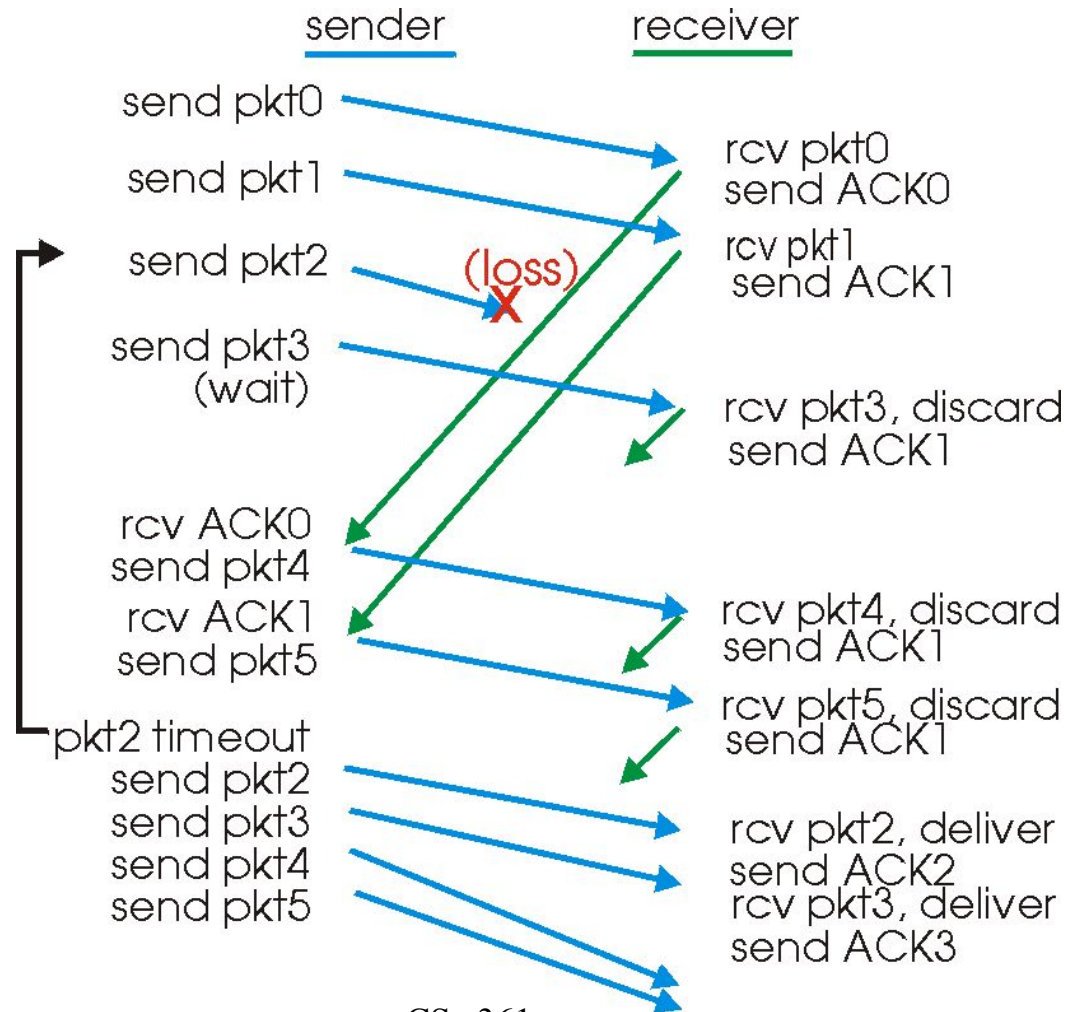
Duplicate acknowledgment

- TCP acknowledgment
 - cumulative acknowledgment
 - example
 - rcv: [0, 500), [500, 1000), [1500, 2000), [2000, 2500)
 - ack: 500, 1000, 1000 (1st dupack), 1000 (2nd dupack)
- Enough duplicate acknowledgments
 - *indicate* packet loss may have occurred
 - ack: 500, 1000, 1000, 1000, 1000 (3rd dupack)
 - packet [1000, 1500) is considered lost

Error recovery

- End-to-end retransmission
 - go-back-N (GBN)
 - retransmit from ackno and upward
 - selective retransmission
 - only retransmit those “known” to be lost
- TCP's error recovery
 - mostly GBN
 - receiver can buffer out-of-order packets
 - explore further: TCP selective acknowledgment

GBN in action ($N = 4$)



This lecture

- TCP error control: wow!
 - purpose (why) and approaches (how)
 - mechanisms (what)
 - detection
 - recovery
- Explore further
 - TCP selective acknowledgment (SACK)
 - <http://www.icir.org/floyd/>

Next lecture

- TCP congestion control
 - * the thing that keeps the Internet still **alive**
 - * Chapter 3, all sections required this month

TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
 - initially `cwnd` = 1 MSS
 - double `cwnd` every RTT
 - done by incrementing `cwnd` for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast

